

Ecommerce Negative Review Agent

A no-code Zapier AI agent that reads customer reviews, classifies sentiment, and drafts an empathetic reply only for the reviews that need one.

Zapier Agents

Google Sheets

LLM System Prompt

No-Code

The Problem

Reading every incoming review doesn't scale — but missing the wrong one does real damage.

01

Not every review needs a reply

Positive and neutral reviews should be left alone — responding to everything wastes effort and reads as robotic.

02

Missing a negative one is costly

A frustrated customer who never hears back can quietly damage trust and brand perception at scale.


03

Manual triage doesn't scale

As review volume grows, a human reading every row to decide "reply or not" becomes the bottleneck.

How It Works

Reviews flow in from Google Sheets; the agent classifies and drafts; only new rows are touched.



Key Decisions & Trade-offs

1

No-code over custom-coded pipeline

Zapier Agents chosen for build speed and zero infra — traded for no version-controlled logic or automated tests.

2

LLM prompt over keyword classifier

One system prompt both classifies sentiment and drafts a context-aware reply, instead of a brittle rules engine.

3

Guardrails iterated through testing

"Don't rewrite the review" proved too narrow; broadened to "don't hallucinate details, customers, actions, or reviews."

4

On Demand, then scheduled hourly

Verified behavior manually first, then switched the trigger to run automatically once trustworthy.

5

Draft only — no auto-send

Gmail is a prerequisite, not a connected tool. A human stays in the loop before anything reaches a customer.

Outcome

What Held Up

- Correctly separated negative reviews from positive/neutral across the test sheet
- Idempotent by design — only touches rows with an empty Mail Draft cell
- Drafts were specific to the issue raised, not generic templates
- No-code build was fast to stand up and easy to hand off

What's Next

- Wire in a Gmail send step with a manual-approval gate, closing the loop end-to-end
- Add an automated test set of reviews to catch prompt regressions before publishing
- Log every agent decision for auditability, not just the final sheet state
- Extend sentiment handling beyond binary reply/no-reply (e.g. flag urgent escalations)