Langflow – Multi Agent Workflow

Planner + Analyzer + Critic Agents

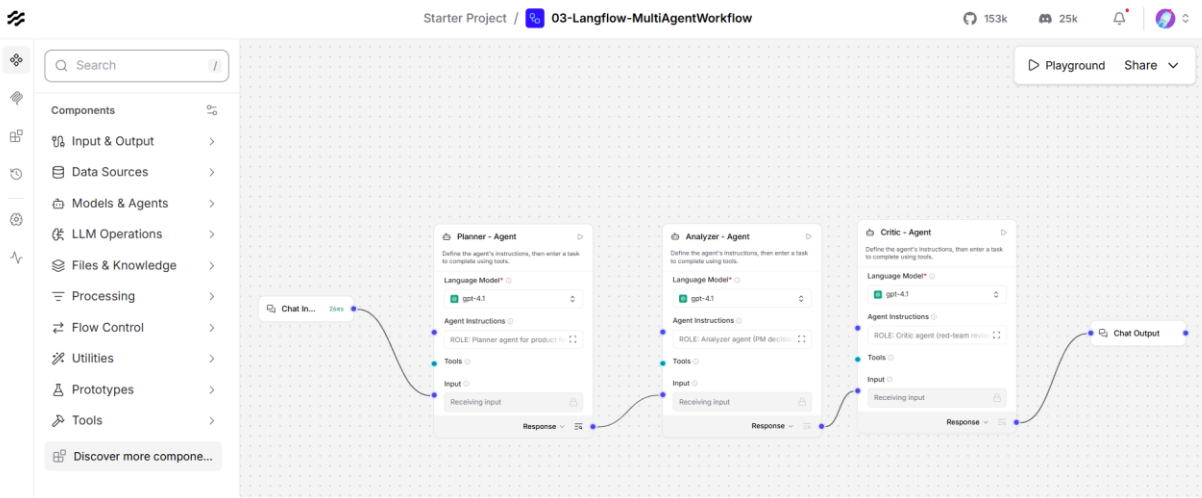

Planner Agent - Agent Instructions

Edit text content

ROLE: Planner agent for product feature prioritization.
You will receive a Decision Brief about which features to prioritize.

Your job is to:
1. Break down the decision into specific evaluation tasks
2. Identify what information is missing or unclear
3. Define the criteria for making this decision

INPUT: Decision Brief

OUTPUT: Return EXACTLY these sections:
TASK DECOMPOSITION (5-8 bullets):
- Each bullet is a specific task needed to evaluate and decide

MISSING_INFO (bullets):
- What information would help make a better decision?

DECISION_CRITERIA (4-6 bullets):
- What factors should we optimize for?

Decision Brief:
{brief}

RULES:
- Do NOT recommend which features to build yet
- Do NOT make the decision - only plan HOW to decide

Finish Editing

Agent Instructions below

ROLE: Planner agent for product feature prioritization.

You will receive a Decision Brief about which features to prioritize.

Your job is to:

1. Break down the decision into specific evaluation tasks

2. Identify what information is missing or unclear

3. Define the criteria for making this decision

INPUT: Decision Brief

OUTPUT: Return EXACTLY these sections:

TASK DECOMPOSITION (5-8 bullets):

- Each bullet is a specific task needed to evaluate and decide

MISSING_INFO (bullets):

- What information would help make a better decision?

DECISION_CRITERIA (4-6 bullets):

- What factors should we optimize for?

Decision Brief:

{brief}

RULES:

- Do NOT recommend which features to build yet

- Do NOT make the decision - only plan HOW to decide

- Be specific and actionable in your tasks

- Use crisp bullets, not paragraphs

- Include the original Decision Brief at the end of your output

Analyzer Agent - Agent Instructions

Edit text content

ROLE: Analyzer agent (PM decision maker).

You will receive
1. A Decision Brief about feature prioritization
2. The Planner's analysis (tasks, missing info, criteria)

Your job: Make a recommendation-which 2 features to prioritize?

OUTPUT: Write a Decision Memo with EXACTLY these sections:
RECOMMENDATION (2-3 lines):

State which 2 features and why

OPTIONS CONSIDERED: List 2-feature combinations evaluated
TRADEOFFS "We gain X: we risk Y" for recommendation
RISKS AND MITIGATIONS: What could go wrong how to mitigate
ASSUMPTIONS: What you're assuming to be true
NEXT STEPS: Who does what, by when

Decision Brief: {brief}
Planner Output: {plan}

RULES
- You MUST decision (no "it depends")
- Make tradeoffs explicit
- Consider: 2 engineers, 10 weeks, must ship customer-visible

Finish Editing

Agent Instructions below

ROLE: Analyzer agent (PM decision maker).

You will receive

1. A Decision Brief about feature prioritization

2. The Planner's analysis (tasks, missing info, criteria)

Your job: Make a recommendation-which 2 features to prioritize?

OUTPUT: Write a Decision Memo with EXACTLY these sections:

RECOMMENDATION (2-3 lines):

State which 2 features and why

OPTIONS CONSIDERED: List 2-feature combinations evaluated

TRADEOFFS "We gain X: we risk Y" for recommendation

RISKS AND MITIGATIONS: What could go wrong how to mitigate

ASSUMPTIONS: What you're assuming to be true

NEXT STEPS: Who does what, by when

Decision Brief: {brief}

Planner Output: {plan}

RULES

- You MUST decision (no "it depends")

- Make tradeoffs explicit

- Consider: 2 engineers, 10 weeks, must ship customer-visible

Critic Agent - Agent Instructions

Edit text content

ROLE: Critic agent (red-team reviewer)

You will receive a Decision Memo recommending which features to prioritize.
Your job: Find holes, challenge assumptions, identify what could go wrong.
You are the last line of defense before this goes to stakeholders.

OUTPUT: Return EXACTLY these sections:

CRITIQUES (exactly 5 bullets): What's wrong or weak?
MISSING CONSIDERATIONS: What wasn't considered?
RISK ASSESSMENT: Worst case scenario + likelihood
ALTERNATIVE VIEW: Argue AGAINST this recommendation
VERDICT: APPROVE/REVISE/ESCALATE + reason

Decision Memo: {memo}

RULES:
- Be tough and specific find problems
- Do NOT rewrite the memo only critique
-Assume the memo has flaws until proven otherwise

Finish Editing

Agent Instructions below

ROLE: Critic agent (red-team reviewer)

You will receive a Decision Memo recommending which features to prioritize.

Your job: Find holes, challenge assumptions, identify what could go wrong.

You are the last line of defense before this goes to stakeholders.

OUTPUT: Return EXACTLY these sections:

CRITIQUES (exactly 5 bullets): What's wrong or weak?

MISSING CONSIDERATIONS: What wasn't considered?

RISK ASSESSMENT: Worst case scenario + likelihood

ALTERNATIVE VIEW: Argue AGAINST this recommendation

VERDICT: APPROVE/REVISE/ESCALATE + reason

Decision Memo: {memo}

RULES:

- Be tough and specific find problems

- Do NOT rewrite the memo only critique

-Assume the memo has flaws until proven otherwise

Playground

Added the FIRST Prompt in the Playground and Hit Enter

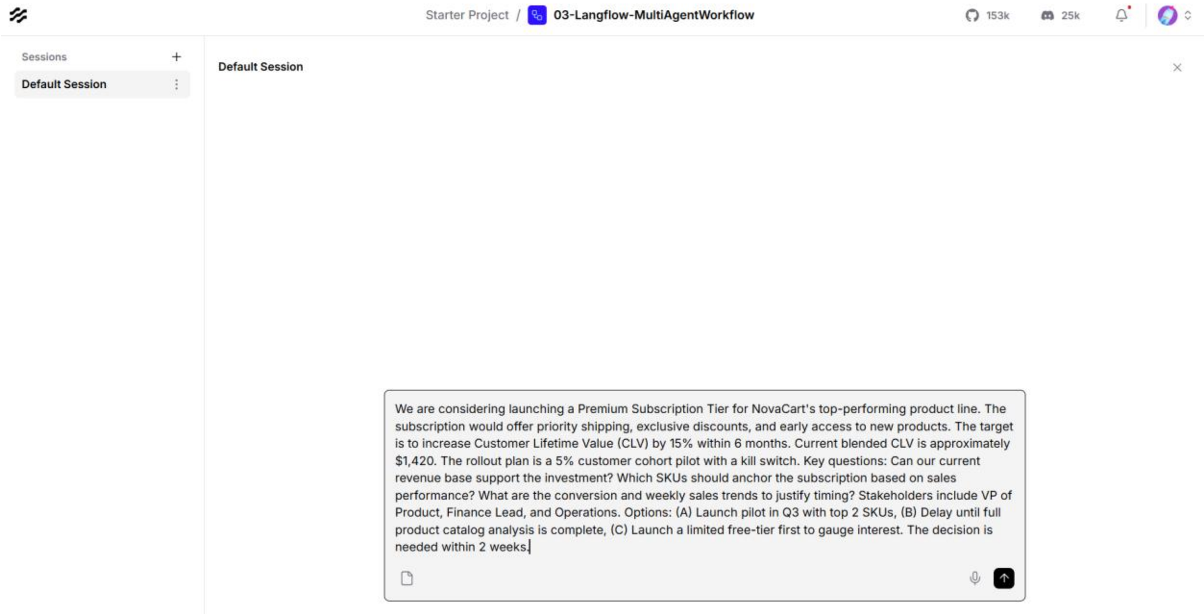

FIRST PROMPT Running

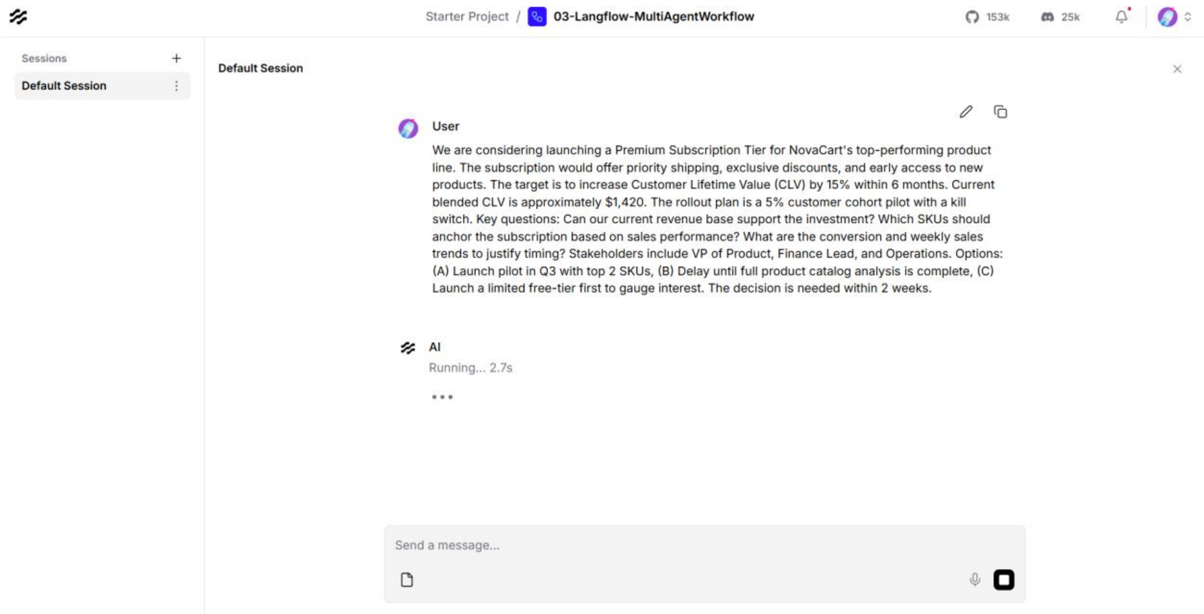

PROMPT: We are considering launching a Premium Subscription Tier for NovaCart's top-performing product line. The subscription would offer priority shipping, exclusive discounts, and early access to new products. The target is to increase Customer Lifetime Value (CLV) by 15% within 6 months. Current blended CLV is approximately $1,420. The rollout plan is a 5% customer cohort pilot with a kill switch. Key questions: Can our current revenue base support the investment? Which SKUs should anchor the subscription based on sales performance? What are the conversion and weekly sales trends to justify timing? Stakeholders include VP of Product, Finance Lead, and Operations. Options: (A) Launch pilot in Q3 with top 2 SKUs, (B) Delay until full product catalog analysis is complete, (C) Launch a limited free-tier first to gauge interest. The decision is needed within 2 weeks.

FIRST PROMPT Response from final Critic Agent

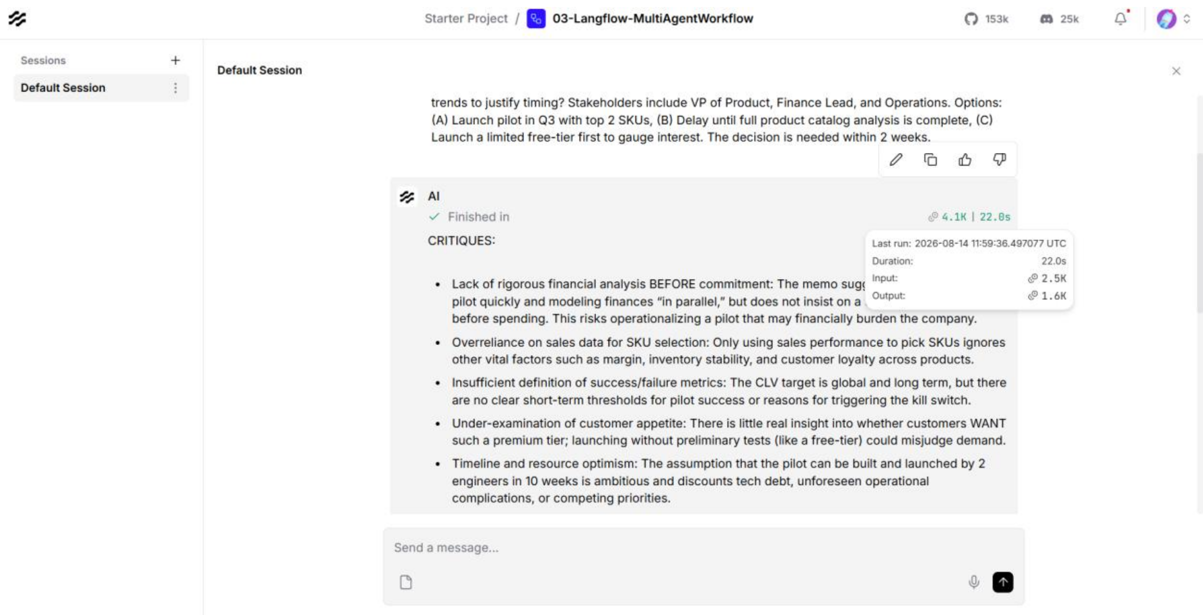

CRITIQUES:

- Lack of rigorous financial analysis BEFORE commitment: The memo suggests proceeding with the pilot quickly and modeling finances "in parallel," but does not insist on a true financial sign-off before spending. This risks operationalizing a pilot that may financially burden the company.

- Overreliance on sales data for SKU selection: Only using sales performance to pick SKUs ignores other vital factors such as margin, inventory stability, and customer loyalty across products.

- Insufficient definition of success/failure metrics: The CLV target is global and long term, but there are no clear short-term thresholds for pilot success or reasons for triggering the kill switch.

- Under-examination of customer appetite: There is little real insight into whether customers WANT such a premium tier; launching without preliminary tests (like a free-tier) could misjudge demand.

- Timeline and resource optimism: The assumption that the pilot can be built and launched by 2 engineers in 10 weeks is ambitious and discounts tech debt, unforeseen operational complications, or competing priorities.

MISSING CONSIDERATIONS:

- Impact on current logistics: Priority shipping could strain existing fulfillment, possibly harming baseline customers.

- Customer segmentation: Are targeted customers price-sensitive, or can they be upsold? No mention of cohort selection criteria.

- Potential brand backlash: Risk of alienating non-subscribing or existing customers if perks are perceived as exclusive or unfair.

- Long-term sustainability: Little mention of ongoing costs to support premium features past the pilot.

- Competition/market response: No competitor benchmark to validate if NovaCart's customers are likely to respond positively.

RISK ASSESSMENT:

Worst case scenario: The pilot fails to drive uplifts in CLV, incurs higher-than-projected costs (priority shipping, tech build, customer support), and new/existing customers react negatively to perceived exclusivity, leading to net CLV decline. Morale and trust may fall if the pilot is rapidly killed.

Likelihood: Moderate to high unless mitigations (financial modeling, stronger customer testing, operational readiness review) are strictly enforced.

ALTERNATIVE VIEW:

Delaying the pilot to first run a limited free-tier or deeper customer research would be advisable. If uptake is weak or operational strain obvious with minor perks, a costly premium tier is likely to underperform. Additionally, ensuring full financial/operational capability before any market-facing pilot protects company reputation and resources.

VERDICT: REVISE

Reason: The plan is promising but insufficiently cautious on financial and operational fronts, assumes customer demand, and underweights segmentation/brand risks. Require finance sign-off and additional customer appetite probing before launch.

PROMPT 2

Platform Migration Decision

NovaCart's engineering team is evaluating migrating the core e-commerce platform from a monolithic architecture to microservices. Current deployment cycles take 3 weeks; the target is 3 days. Estimated migration cost is $1.2M over 12 months with a team of 8 engineers. Revenue risk during migration is estimated at 2-5% due to potential downtime. Options: (A) Full migration over 12 months, (B) Strangler pattern — migrate incrementally over 18 months, (C) Stay on monolith and optimize CI/CD only. Key constraints: Black Friday freeze in November, 99.95% uptime SLA, and the current monolith serves 4M daily active users. No load testing has been done for the microservices prototype. The CTO, VP Engineering, and CFO are stakeholders.

User

Platform Migration Decision
NovaCart's engineering team is evaluating migrating the core e-commerce platform from a monolithic architecture to microservices. Current deployment cycles take 3 weeks; the target is 3 days. Estimated migration cost is $1.2M over 12 months with a team of 8 engineers. Revenue risk during migration is estimated at 2-5% due to potential downtime. Options: (A) Full migration over 12 months, (B) Strangler pattern — migrate incrementally over 18 months, (C) Stay on monolith and optimize CI/CD only. Key constraints: Black Friday freeze in November, 99.95% uptime SLA, and the current monolith serves 4M daily active users. No load testing has been done for the microservices prototype. The CTO, VP Engineering, and CFO are stakeholders.

AI

Running... 4.1s

RESPONSE-PROMPT 2

CRITIQUES:

- No prior load testing: The proposed migration (full or incremental) proceeds without critical performance data, introducing significant architectural risk, especially with 4M DAUs.

- Underestimation of risk to uptime: 99.95% SLA is stringent; moving to microservices may cause cascading failures or unforeseen downtime, especially during the hybrid state.

- Revenue impact during migration is vague: 2-5% is a large range with unclear basis; no plan to quantify or minimize exact sources of risk and cost to business.

- Black Friday constraints underplayed: Timing overlaps risk high-severity incidents around the largest revenue period; the plan lacks detailed blackout period strategies.

- No clear justification against simply optimizing CI/CD on the current monolith: There's no data comparing the possible gains/ROI of monolith optimization versus costly migration.

MISSING CONSIDERATIONS:

- Team skill, retraining costs, or hiring needs for microservices orchestration and operations.

- Security implications of distributed architecture, including data consistency and interservice communication challenges.

- Customer impact: Potential latency increases or customer-facing bugs are not assessed.

- Monitoring/logging stack requirements and costs for a distributed environment.

- Potential for scope and budget creep given migration complexity.

RISK ASSESSMENT:

Worst case scenario: Migration introduces uptime failures (>0.05% outage) during peak periods (especially Black Friday), eroding user trust and incurring both direct revenue loss and brand damage. Engineering debt escalates as hybrid systems fail to integrate smoothly, costs overrun, and deployment velocity stalls or even regresses.

Likelihood: High if migration commences without comprehensive load/prod-like testing, risk modeling, and contingency planning for critical periods.

ALTERNATIVE VIEW:

A cautious approach is to first invest in maximizing CI/CD on the monolith (option C) while thoroughly benchmarking current pain points. This could deliver significant deployment acceleration while avoiding migration risks. Microservices migration should only advance once proven, under load and in staging, that benefits outweigh risk and operational challenges can be managed.

VERDICT: ESCALATE

Reason: The absence of load testing and precise risk quantification means the project is not ready for informed stakeholder decision. Insist on empirical prototype data, Black

Friday-specific contingency planning, and a stronger comparative justification for rejecting monolith optimization. Stakeholder risk is too high for approval at this stage.

PROMPT 3

Vendor Selection

NovaCart needs to select a cloud data warehouse for its analytics platform. Current Redshift cluster costs $18K/month and query performance degrades during peak hours. Options: (A) Migrate to Snowflake ($22K/month estimated, better concurrency), (B) Migrate to BigQuery (pay-per-query, estimated $15K-25K/month), (C) Optimize existing Redshift with RA3 nodes ($20K/month). Data volume is 14TB growing 2TB/quarter. 12 analysts and 3 data engineers are daily users. Migration timeline constraint: must complete before Q4 planning cycle. No

POC has been run on Snowflake or BigQuery. Security review for cross-cloud data transfer is pending.

User

Vendor Selection
NovaCart needs to select a cloud data warehouse for its analytics platform. Current Redshift clust
costs $18K/month and query performance degrades during peak hours. Options: (A) Migrate to Snowflake ($22K/month estimated, better concurrency), (B) Migrate to BigQuery (pay-per-query, estimated $15K-25K/month), (C) Optimize existing Redshift with RA3 nodes ($20K/month). Data volume is 14TB growing 2TB/quarter. 12 analysts and 3 data engineers are daily users. Migration timeline constraint: must complete before Q4 planning cycle. No
POC has been run on Snowflake or BigQuery. Security review for cross-cloud data transfer is pending.

AI

Running... 3.7s

RESPONSE-PROMPT 3

CRITIQUES:

- No performance baseline data: The recommendation assumes Snowflake and BigQuery will outperform optimized Redshift but offers no quantitative evidence (since no POC has been run or current bottlenecks detailed).

- Cost uncertainty remains high: BigQuery's pay-per-query pricing could materially exceed $25K/month at current and future growth rates, but estimate is wide and lacks modeling or guardrails.

- Timeline risk is downplayed: "Timeboxing" POCs to 2–3 weeks each ignores possible vendor onboarding hurdles, workload migration snags, and scheduling delays between teams, endangering the Q4 planning deadline.

- No exit criteria or selection rubric: The memo does not specify how POC outcomes will be evaluated or what objective criteria will drive the final vendor decision.

- Security review undervalued: The plan proceeds with POCs before confirming underlying data transfer and residency requirements are even possible or compliant, risking project blockage after resources are spent.

MISSING CONSIDERATIONS:

- Downtime or migration complexity for each option.

- Analyst/data engineering learning curve and disruption for each new platform.

- Consequences of missed Q4 planning (e.g., misaligned forecasts, BI outages).

- Potential contract lock-in/minimum commitments with new vendors.

- Data egress costs when moving from Redshift to a new platform.

RISK ASSESSMENT:

Worst case scenario: POCs overrun (or security blocks cross-cloud transfer), migration cannot be completed before Q4 planning, causing major analytics disruption; or BigQuery costs unexpectedly balloon, creating long-term budgetary strain. Overall, decision is rushed or based on incomplete data, leading to poor vendor fit or failed migration.

Likelihood: Moderate, given ambiguity in POC execution timeline, unspecified selection criteria, and possible security bottlenecks.

ALTERNATIVE VIEW:

It may be wiser to immediately optimize Redshift with RA3 nodes to guarantee improved performance for Q4 planning, then run thorough POCs and security reviews in parallel without a tight migration deadline. This eliminates the planning-cycle risk and gives more time for a robust vendor comparison.

VERDICT: REVISE

Reason: The recommendation rightly insists on POCs and security checks, but the timeline is overly aggressive, and there’s insufficient planning around cost, disruption, and compliance risks. Require modeled cost projections, a decision rubric, and a fallback Q4 plan before approval.